

80.6x drop-in on B300; up to 775x with compact L2 residency

The client-supplied benchmark isolates a real performance issue: memory geometry, data movement, and math pipeline choice. The report keeps the safe drop-in answer separate from compact-data L2 residency paths.

ORIGINAL BASELINE

14.8900 ms

Strict problem definition

RECOMMENDED DROP-IN

80.6x

0.1848 ms

CURRENT COMPACT L2

638.7x

0.0233 ms

FASTEST OBSERVED L2

774.9x

0.0169 ms on RTX PRO 6000
Blackwell, 1 GPU

BEST SCORE PATH

595.9x

0.0250 ms

Executive Summary

The supplied problem is a strong diagnostic: it processes a fixed 8192 x 8192 float grid with five dependent transcendental iterations per element and a 1e-3 relative tolerance check.

The recommended production default is the vectorized float kernel. It keeps the input/output contract intact and moves the runtime from 14.8900 ms to 0.1848 ms on NVIDIA B300 SXM6 AC.

The compact L2 path is the strategic upside. On this run it reaches 0.0233 ms, or 638.7x; the fastest observed compact L2 row across the report set is 0.0169 ms / 774.9x on RTX PRO 6000 Blackwell, 1 GPU. These rows are specialized contract results, not drop-in replacements.

Nsight classifies the drop-in float path as compute-throughput sensitive on the measured GPU, which explains why fixed-range affine math can win materially here. The compact fused paths still matter because they remove global-memory traffic and launch boundaries.

Recommendation

Adopt the vectorized float kernel as the production default.

Treat compact U8 and L2-resident paths as second-phase options for systems that own the producer/consumer interface. They are compelling when surrounding stages can store or consume compact data directly and avoid avoidable HBM round trips.

Adoption Path

The benchmark results separate the immediate production answer from higher-upside paths that require more control over the surrounding dataflow.


PATH	REPRESENTATIVE RESULT	CUSTOMER ACTION
Adopt now	0.185 ms / 80.6x	Same float input/output contract. This is the production recommendation when the surrounding application must stay unchanged.
Narrow the data contract	0.0321 ms / 463.9x	Use when adjacent stages can produce or consume compact U8 values directly instead of expanding every lane back to floats.
Use compact L2 residency	0.0169 ms / 774.9x RTX PRO 6000 Blackwell, 1 GPU	Best for latency-sensitive producer/consumer systems that can keep compact working data local. It is HBM-light, not HBM-free, and should be remeasured on each target SKU.
Fuse the downstream score	0.0250 ms / 595.9x	Use when the next stage needs the score stream, not the intermediate compact U8 output buffer.

Client Options

DROP-IN Vectorized float default Recommended when the float input/output contract must stay unchanged. 80.6x 0.1848 ms 512 MiB data moved	OUTPUT CONTRACT FP16 output Useful when the client can accept half-precision output storage. 236.7x 0.0629 ms 256 MiB data moved	COMPACT CONTRACT Compact U8 input + FP16 output Best when an upstream producer can emit only the two changing values directly. 374.1x 0.0398 ms 160 MiB data moved
SPECIALIZED CONTRACT uint4 + L2 U8 in/out kernel Fastest kernel-only result: compact U8 input/output with persisting input L2. 638.7x 0.0233 ms 64 MiB data moved	SPECIALIZED PIPELINE uint4 + L2 compact pipeline Best adjacent producer-consumer path: uint4 producer plus persisting compact output. 325.0x 0.0458 ms 160 MiB data moved	FUSED PIPELINE Fused compact U8 score + L2 Fastest score-producing path if compact input is reused and can stay resident in L2. 595.9x 0.0250 ms 96 MiB data moved

Performance Ladder


Client-relevant variants, sorted by speedup against the original client baseline. L2-resident results are included when they were measured in the current report.



Why The Speedup Holds


Data movement falls as the contract narrows

The drop-in path fixes memory access geometry. Data-contract-changing paths then reduce output bytes and keep only the two lanes that actually vary in this benchmark.



Profiler evidence

The raw L1 sector count is a transaction counter: it confirms that vectorized loads and stores are coalesced into predictable memory chunks. The client takeaway is bandwidth pressure, not a need to inspect sector math.



Ranked Benchmark Portfolio

Sorted by speedup vs the original client baseline, including L2-resident rows when available. Footprint combines data moved, register pressure, and spill status so the table stays decision-oriented.

VARIANT	TRACK	STATUS	SPEEDUP	TIME	FOOTPRINT
uint4 + L2 U8 in/out kernel Fastest kernel-only result: compact U8 input/output with persisting input L2.	SPECIALIZED CONTRACT	Pass	638.7x	0.0233 ms	64 MiB data moved
uint4-packed U8 in/out kernel Processes eight compact groups per thread; faster producer without relying on L2 hints.	SPECIALIZED CONTRACT	Pass	632.2x	0.0236 ms	64 MiB data moved
Fused compact U8 score + L2 Fastest score-producing path if compact input is reused and can stay resident in L2.	FUSED PIPELINE	Pass	595.9x	0.0250 ms	96 MiB data moved
L2-persisting U8 in/out kernel Fastest kernel-only result when compact input is reused and fits in persisting L2.	SPECIALIZED CONTRACT	Pass	540.5x	0.0276 ms	64 MiB data moved
Fused compact U8 input to score Reads compact U8 input and writes score output directly, skipping the compact output buffer.	FUSED PIPELINE	Pass	470.6x	0.0316 ms	96 MiB data moved
Compact U8 input + U8 output Fastest kernel-only result; does not include a downstream consumer.	SPECIALIZED CONTRACT	Pass	463.9x	0.0321 ms	64 MiB data moved 14 registers/thread 0 spill bytes
Compact U8 input + FP16 output Best when an upstream producer can emit only the two changing values directly.	COMPACT CONTRACT	Pass	374.1x	0.0398 ms	160 MiB data moved 28 registers/thread 0 spill bytes
Compact U16 input + FP16 output Preserves tolerance while storing only the two changing values in each four-value group.	COMPACT CONTRACT	Pass	352.8x	0.0422 ms	192 MiB data moved 30 registers/thread 0 spill bytes
uint4 + L2 compact pipeline Best adjacent producer-consumer path: uint4 producer plus persisting compact output.	SPECIALIZED PIPELINE	Pass	325.0x	0.0458 ms	160 MiB data moved
Compact FP32 two-value input + FP16 output Upper-bound compact input shape before fixed-point quantization.	COMPACT CONTRACT	Pass	312.2x	0.0477 ms	256 MiB data moved 30 registers/thread 0 spill bytes
uint4 U8 producer + consumer Vectorized compact producer followed by the existing compact consumer.	PIPELINE	Pass	288.3x	0.0517 ms	160 MiB data moved
L2-resident U8 in/out pipeline U8 input, U8 output, and compact consumer timed together with persisting L2.	SPECIALIZED PIPELINE	Pass	238.2x	0.0625 ms	160 MiB data moved
FP16 output Useful when the client can accept half-precision output storage.	OUTPUT CONTRACT	Pass	236.7x	0.0629 ms	256 MiB data moved 30 registers/thread 0 spill bytes

VARIANT	TRACK	STATUS	SPEEDUP	TIME	FOOTPRINT
Affine float default Documents the fixed-range math opportunity, but does not materially beat the float default.	DROP-IN-SPECIALIZED	Pass	176.0x	0.0846 ms	512 MiB data moved 22 registers/thread 0 spill bytes
Compact U8 output + score pipeline Setup-paid compact-output pipeline; retained as a comparison, not the extreme path.	PIPELINE	Pass	128.1x	0.1162 ms	448 MiB data moved
Vectorized float default Recommended when the float input/output contract must stay unchanged.	DROP-IN	Pass	80.6x	0.1848 ms	512 MiB data moved 26 registers/thread 0 spill bytes
Scalar coalesced kernel Shows the value of coalesced memory access before vectorization.	STRUCTURAL FIX	Pass	20.7x	0.7197 ms	512 MiB data moved 21 registers/thread 0 spill bytes
Original row-stride kernel Same harness, inefficient memory layout	KERNEL BASELINE	Pass	2.3x	6.4807 ms	512 MiB data moved 19 registers/thread 0 spill bytes

Maximum-Performance L2 Path

When a compact data contract is worth considering

This path is aimed at latency-sensitive customers who can own the compact data contract, producer/consumer coupling, and maintenance burden. It is not the safe library default.

There are two useful combinations. First, the current best kernel-only shape uses uint4 loads/stores to process eight compact groups per thread, then keeps compact U8 input in persisting L2; that measured 0.0233 ms, or 638.7x against the original client baseline. Second, the uint4 compact producer can write U8 x/w output and keep that output resident for an adjacent compact consumer; that producer-plus-consumer path measured 0.0458 ms, or 325.0x.

The kernel-only L2 result is the fastest measured row on this system, but it depends on compact input reuse and a working set that fits the persisting-L2 budget. On B200-class systems, very high HBM3e bandwidth narrows the cache advantage, so this tactic should be remeasured rather than assumed.

This is an HBM-light path rather than an HBM-free one. It removes avoidable round trips and reduces read pressure, but the compact consumer still shows a memory-path profile because it writes the score stream and moves predictable global-memory transactions.

COMPACT OUTPUT FOOTPRINT 32 MiB	MEASURED L2 CACHE 126 MiB	PERSISTING-L2 BUDGET 79 MiB
KERNEL-ONLY L2 RESULT 0.0233 ms	L2 LIFT ON BEST PRODUCER 1.01x	UINT4 PRODUCER LIFT 1.18x
SCALAR L2 VS THRASH 1.55x	WARM VS COLD CONSUMER 1.3x	PERSISTING TOTAL LIFT 1.13x

UINT4 PIPELINE LIFT

1.36x

FUSED SCORE RESULT

0.0250 ms

FUSED VS BEST ADJACENT

1.83x

FUSED VS STRICT BASELINE

595.9x

Measured latency path

Producer after L2 thrash

Current best kernel shape after cache pressure



0.0427 ms

Producer warm

Current best kernel shape with natural reuse



0.0319 ms

Producer with persisting input


Current best kernel shape plus L2 input residency



0.0276 ms

uint4 producer

Eight compact groups per thread



0.0236 ms

uint4 + L2 producer

Best kernel-only shape



0.0233 ms

Thrashed consumer

Cold path after a 256 MiB L2-thrashing pass



0.0326 ms

Warm consumer

Producer output is still resident



0.0261 ms

Persisting consumer

CUDA persisting-L2 window applied



0.0223 ms

Fused score + L2

Direct compact input to score with persisting compact input



0.0250 ms

Fused score

Direct compact input to score, no intermediate output buffer



0.0316 ms

Persisting total

Producer plus compact consumer



0.0625 ms

Cache residency planning model

Planning estimate only: assume about 70% of advertised cache is usable for resident working data, then verify on the target SKU. Aggregate cache helps only when the workload is partitioned so each GPU or die keeps its shard local.

GPU FAMILY	USABLE CACHE ESTIMATE	256 MIB WORKING SET	512 MIB WORKING SET	READOUT
H100 / H200	~35 MiB	8 GPUs	15+ GPUs	L2 density is too small for this tactic at practical scale.

GPU FAMILY	USABLE CACHE ESTIMATE	256 MIB WORKING SET	512 MIB WORKING SET	READOUT
RTX 6000 Ada / RTX 5090	~67 MiB	4 GPUs	8 GPUs	Good single-die cache, but still needs sharding for larger working sets.
RTX Pro 6000 Blackwell	~90 MiB	3 GPUs	6 GPUs	Best fit among workstation-class NVIDIA parts in this planning model.
B200	~180 MiB logical	2 GPUs	3 GPUs	High HBM3e bandwidth narrows the cache advantage; use L2 mainly for reuse and latency.
B300 SXM6 AC (measured)	~89 MiB	3 GPUs	6 GPUs	Measured in this B300 run: 126.5 MiB L2 and 79.1 MiB persisting budget.

Hardware Key Result Comparison

Representative rows from earlier L4 and RTX PRO 6000 runs are shown alongside the current measured system. The RTX PRO 6000 two-GPU entry is a topology note only; no multi-GPU timing is claimed.

HARDWARE	REPRESENTATIVE OPTION	ROLE	RESULT	READOUT
NVIDIA L4	Vectorized float default	Safe drop-in	2.31 ms / 15.1x	Earlier L4 campaign measured the float-compatible optimization track; compact and L2 variants were not part of that run.
RTX PRO 6000 Blackwell, 1 GPU	Vectorized float default	Safe drop-in	0.310 ms / 42.2x	Same float input/output contract as the baseline benchmark.
RTX PRO 6000 Blackwell, 1 GPU	uint4 + L2 U8 in/out kernel	Specialized kernel-only	0.0169 ms / 774.9x	This is the maximum-performance kernel-only option: compact U8 input/output with persisting input L2. It is not a drop-in float path and does not include downstream score work.
RTX PRO 6000 Blackwell, 1 GPU	uint4 + L2 compact pipeline	Specialized producer + consumer	0.0419 ms / 312.1x	Adjacent compact producer and compact consumer timed together; useful for coupled custom-data pipelines.
RTX PRO 6000 Blackwell, 2-GPU host	Topology note	Hardware context	PHB path, no NVLink timing claim	The earlier two-GPU host showed PCIe/PHB connectivity. No multi-GPU timing was committed, so this is not a performance row.
Current run: NVIDIA B300 SXM6 AC	Vectorized float default	Safe drop-in	0.185 ms / 80.6x	Recommended production default when the float input/output contract must remain unchanged.
Current run: NVIDIA B300 SXM6 AC	uint4 + L2 U8 in/out kernel	Specialized kernel-only	0.0233 ms / 638.7x	Fastest kernel-only row in the current measurement set; it requires the compact data contract.
Current run: NVIDIA B300 SXM6 AC	Fused compact U8 score + L2	Specialized score path	0.0250 ms / 595.9x	Best score-producing row in the current report; it is separate from the kernel-only producer result.

Multi-GPU Outlook

Single-GPU host

This run does not contain a multi-GPU opportunity because only one GPU is visible.

ONE-GPU HARNESS
MEMORY
784.00 MiB

PARTITIONED HARNESS
FITS
yes

The current 8192 x 8192 harness fits on one GPU with headroom, so the next multi-GPU benchmark should be a row-sharded throughput test, not a replacement for the single-GPU score above.

GPU-TO-GPU PATH
unknown

NVLINK
no

Run Context

ITEM	VALUE
GPU	NVIDIA B300 SXM6 AC
GPU count	1
GPU interconnect	unknown (no NVLink)
Driver	580.126.09
CUDA compiler	Cuda compilation tools, release 13.0, V13.0.88
Problem size	8192 x 8192

Source Package Links

These relative links are intended for the source bundle sent to the client, so the report, README, and implementation stay connected when reviewed offline.

SOURCE	PURPOSE
README	Project overview, hardware scoreboards, and experiment ledger.
CUDA source	Optimized kernel, affine/compact variants, and L2 experiment harness.
Report renderer	HTML report generator used for this PDF and published page.